HW1: Sentiment Classification

Jonah Philion
jonahphilion@college.harvard.edu
github

February 2, 2018

1 Introduction

In this problem set, we attempt to classify movie reviews as positive or negative. All models investigated take the form of

$$p(y_i) = \sigma(W\phi(x) + b)$$

where $p(y_i)$ is the probability a review x is negative. The models studied are

- naive bayes
- logistic regression with "bag of words" features
- multi-layer perceptron with "continuous bag of words" features
- convolutional neural net

And just for fun

- Fine-tuning ResNet classifier on images of text

2 Problem Description

For all models, a sentence $\boldsymbol{x}_i$ is encoded as a sequence $x_1, ..., x_n$ where each x_j is a one-hot vector of length the vocabulary $\mathcal{V}$. The classification y_i associated with $\boldsymbol{x}_i$ is 0 if $\boldsymbol{x}_i$ is positive and 1 if $\boldsymbol{x}_i$ is negative. Embeddings $\mathcal{E}$ map a one hot vector x_j to a dense vector of size d.

3 Model and Algorithms

All models are trained on the Stanford Sentiment Treebank (SST1). Unless otherwise specified, models requiring gradient descent are trained with an Adam optimizer of learning rate 0.001 and weight decay 0.0005. Training loss and validation loss are recorded in real time with visdom.

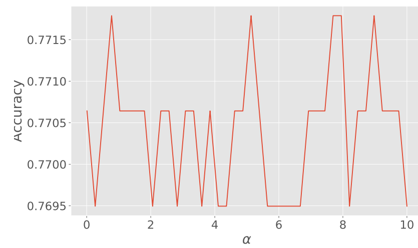

Figure 1: Affect of global smoothing parameter α on validation accuracy. Validation accuracy is calculated using sklearn's `accuracy_score`. The performance is not sensitive to α.

4 Experiments

4.1 Naive Bayes

All words for both labels are initialized to a count of α. Figure 1 shows that performance of the algorithm on the validation set is robust to changes in α. With $\alpha = .5$, performance on the test set using binarized and count word vectors are analyzed in Table 1.

The weight vector determined by Naive Bayes can be interpreted as the sentiment associated with particular words. The most positive and most negative words and their weights for binarized bag of words are shown in Table **??**.

Model	Acc.	Bce.	Roc.
BINARIZED	0.791	0.679	0.867
COUNTS	0.793	0.686	0.867

Table 1: Binarizing or counting words does not significantly affect the performance of the model on the test set.

stupid	unfunny	pointless	poorly	suffers	Feels	tiresome	car
5.39452	5.1085	5.03998	4.96641	4.96641	4.88701	4.88701	4.88701
powerful	solid	perfectly	inventive	refreshing	riveting	wonderfully	universal
-5.79766	-5.7109	-4.99019	-4.85754	-4.78398	-4.70458	-4.61832	-4.61832

Table 2: Words with the highest and lowest weights in the naive bayes weight vector agree with intuition for being negative and positive words respectively.

4.2 Logistic Regression

In the logistic regression model, we take the bag of words $\phi(x)$ but train w and b instead of calculating w and b. Figure **??** shows that as the model decreases cross entropy, the accuracy on the validation set also increases. However, this model does not do significantly better than naive bayes.

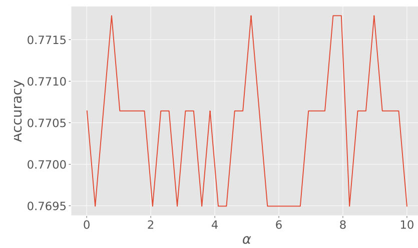

Figure 2: The model does not overfit.

4.3 Continuous bag of Words

4.4 CNN

4.5 ResNet

5 Conclusion

End the write-up with a very short recap of the main experiments and the main results. Describe any challenges you may have faced, and what could have been improved in the model.