

HW2: Language Modeling

Jonah Philion
jonahphilion@college.harvard.edu
github

February 17, 2018

1 Introduction

In this problem set, we attempt to classify movie reviews as positive or negative. All models investigated take the form of

$$x_{i+1} \sim \sigma(W\phi(x_1, \dots, x_i) + b)$$

where x_1, \dots, x_{i+1} are word vectors for a sentence and σ is the softmax function. The models I try are

- trigram with linear interpolation
- NN language model
- LSTM language model

And just for fun

- Embedding \rightarrow Embedding
- Fine-tuning ResNet classifier on images of text

2 Problem Description

Given 10 consecutive words, we would like to predict the next word. For all models, a sentence x_i is encoded as a sequence x_1, \dots, x_n where each x_j is a one-hot vector of length the vocabulary \mathcal{V} . The model outputs a categorical distribution over the vocabulary \mathcal{V} . Embeddings \mathcal{E}_d map a one hot vector x_j to a dense vector of size d .

3 Model and Algorithms

All models are trained on the Penn Treebank. For models requiring gradient descent, training loss and validation loss are recorded in real time with visdom. Final Mean Average Precision (MAP) is calculated on a validation set.

3.1 Evaluation

1. All models have a predict function which given a batch of sentence fragments outputs an array containing the probabilities for the next word for all batches.
2. All models are tested with the same evaluation function. The function evaluates the model on batches of size 67 of bptt_len 11 of “valid.txt”. The first 10 words of the batch are fed into the predict function and the 11th word is treated as ground truth. The top 20 words are then measured against the correct word with MAP.

The validation is assumed to be fairly similar to the Kaggle set as justified by the fact that the MAP-20 for the trigram model on the validation set is 0.296 and on Kaggle the model receives an MAP of 0.292. Unless explicitly mentioned, all MAP reported below are calculated with this evaluation code, not on Kaggle.

4 Experiments

4.1 Trigram

The interpolated trigram model states

$$p(y_{i+1}) = \alpha \cdot [p_{y_{i-1}, y_i}^{y_{i+1}}, p_{y_i}^{y_{i+1}}, p^{y_{i+1}}]$$

Where α is a vector of length 3 with $|\alpha|_1 = 1$, and p_b^a is the probability of word a given sequence b as determined by count data in the training set.

The unary probabilities are implemented as a vector of size \mathcal{V} , the binary probabilities are $\mathcal{V} \times \mathcal{V}$, and the ternary probabilities are a dictionary with keys (w_1, w_2) and values vectors of size \mathcal{V} .

During training, all sequences of size 1,2 and 3 are counted. The counts are then normalized over \mathcal{V} . At test time, we for the most part return $\alpha \cdot p$. However, for sequences of length 1 not seen in training, we replace p_{binary} with p_{unary} and for sequences of length 2 not seen in training, we check if p_{binary} is in training and if so use p_{binary} , if not use p_{unary} .

To determine the best alpha, we sample randomly from vectors of length 3 that sum to 1. We do so 5000 times and use the best one. The alpha chosen is

$$\alpha = [0.4306712668382596, 0.4897915705677378, 0.07953716259400256]$$

Scores for the above and other alpha of interest are reported below.

4.2 NN Language Model

4.3 LSTM language model

4.4 Embedding \rightarrow Embedding

4.5 ResNet

5 Conclusion

Four standard NLP models and a fifth model which has no predefined understanding of a vocabulary are trained and evaluated on SST1. The best results from all models are shown in Figure 1.

Model	Acc.	Bce.	Roc.
NB Counts	0.793	0.686	0.867
Binarized LR	0.789	0.496	0.876
Mean CBOW	0.824	0.416	0.898
GloVe CNN	0.833	0.577	0.907
ResNet18	0.774	0.685	0.855

Table 1